

# Boosting and Support Vector Machines

By Adrian Chavez-Loya

You're working for a car manufacturer that is looking to implement driver assistance features such as automated steering and adaptive cruise control. While technologically advanced, these systems still require driver attention. Some manufacturers simply require keeping your hands on the wheel but your company would also like to ensure the driver's focus remains on the road. To accomplish this, they'd like you to construct a model that can use the position of facial features to determine whether the driver is looking straight or not.

A separate system has been used to extract the eye, mouth, and nose positions from images taken of the driver, your goal is to use these features to predict the direction of the driver's gaze. The dataset listed below has been provided for these tasks.

## Relevant Dataset

`drivPoints.txt`

- Response Variable: `label` . Note: this includes looking left, right, and straight. We will convert this to a binary response.
- Predictor Variables:
  - `[ xF yF wF hF ]` = face position
  - `[ xRE yRE ]` = right eye position
  - `[ xLE yL ]` = left eye position
  - `[ xN yN ]` = Nose position
  - `[ xRM yRM ]` = right corner of mouth
  - `[ xLM yLM ]` = left corner of mouth

## Source

<https://archive.ics.uci.edu/ml/datasets/DrivFace>

## Task 1: Import the dataset and create a binary variable of `lookingStraight`. Split into train/test set.

This variable should take the value of `1` when `label=2` and `0` everywhere else. There should be a large class imbalance between looking straight or not (which you would expect given the people are driving).

```
In [ ]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split

df = pd.read_csv('drivPoints.txt')
df.head()
```

```
Out [ ]: 
```

	fileName	subject	imgNum	label	ang	xF	yF	wF	hF
0	20130529_01_Driv_001_f	1	1	2	0	292	209	100	112
1	20130529_01_Driv_002_f	1	2	2	0	286	200	109	128
2	20130529_01_Driv_003_f	1	3	2	0	290	204	105	121
3	20130529_01_Driv_004_f	1	4	2	0	287	202	112	118
4	20130529_01_Driv_005_f	1	5	2	0	290	193	104	119

```
In [ ]: # Creted binary response variable 'lookingStraight'

df['lookingStraight'] = np.where(df['label'] == 2, 1, 0)
```

```
In [ ]: ## Split variables (into features and target)
X = df[['xF', 'yF', 'wF', 'hF', 'xRE', 'yRE', 'xLE', 'yLE', 'xN', 'yN', 'wN', 'hN']]
y = df['lookingStraight']
```

```
In [ ]: # Split into subsets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_si
```

## Task 2: Perform a cross-validated (or use a single validation set) grid search of the

## hyperparameters for the GradientBoostingClassifier to find the best model.

You should at least tune the learning rate and number of trees in the model but feel free to go as deep as you'd like on this analysis).

```
In [ ]: from sklearn.ensemble import GradientBoostingClassifier
        from sklearn.model_selection import GridSearchCV

        # Parameter grid defined
        param_grid = {
            'learning_rate': [0.01, 0.1, 0.2, 0.3],
            'n_estimators': [100, 200, 300],
            'max_depth': [3, 5, 7]
        }

        # Created classifier
        gbc = GradientBoostingClassifier()

        # CV with grid search
        grid_search_gbc = GridSearchCV(estimator=gbc, param_grid=param_gr
        grid_search_gbc.fit(X_train, y_train)

        best_params_gbc = grid_search_gbc.best_params_
        best_score_gbc = grid_search_gbc.best_score_

        best_params_gbc, best_score_gbc

Out[ ]: ({'learning_rate': 0.3, 'max_depth': 5, 'n_estimators': 300},
        0.9607603092783507)
```

- We got an accuracy score of 96!
- Hyperparameters are as follows:
  1. learning\_rate : 0.1
  2. max\_depth : 5
  3. n\_estimators : 200

**Task 3: Perform a cross-validated (or use a single validation set) grid search of the hyperparameters for the SVC (Support Vector Classifier) to find the best model.**

You should at least tune `C` and the `kernel` but feel free to go as deep as you'd like on this analysis).

```
In [ ]: from sklearn.svm import SVC

# Parameter grid for SVC
param_grid_svc = {
    'C': [0.1, 1, 10, 100],
    'kernel': ['linear', 'rbf', 'poly'],
    'gamma': ['scale', 'auto']
}
```

```
In [ ]: svc = SVC() #SVC

# CV grid search
grid_search_svc = GridSearchCV(estimator=svc, param_grid=param_gr
grid_search_svc.fit(X_train, y_train)
best_params_svc = grid_search_svc.best_params_
best_score_svc = grid_search_svc.best_score_
best_params_svc, best_score_svc
```

## Testing F1 Scores for both models to test for performance

```
In [ ]: from sklearn.metrics import f1_score
best_gbc = GradientBoostingClassifier(learning_rate=0.1, max_dept
best_gbc.fit(X_train, y_train)
y_pred_gbc = best_gbc.predict(X_test)

best_svc = SVC(C=best_params_svc['C'], kernel=best_params_svc['ke
best_svc.fit(X_train, y_train)
y_pred_svc = best_svc.predict(X_test)

# F1 scores for both models
f1_gbc = f1_score(y_test, y_pred_gbc)
f1_svc = f1_score(y_test, y_pred_svc)

f1_gbc, f1_svc
```

- Looks like my computer is taking forever to everything in task 3! I will make some adjustments to make it easier to run

# Using Random Search CV to redefine parameter grid (full code)

```
In [ ]: from sklearn.model_selection import train_test_split, RandomizedS
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.svm import SVC
from sklearn.metrics import f1_score

# Load the dataset
df = pd.read_csv('drivPoints.txt')

# Create the binary response variable
df['lookingStraight'] = np.where(df['label'] == 2, 1, 0)

# Split into features and target
X = df[['xF', 'yF', 'wF', 'hF', 'xRE', 'yRE', 'xLE', 'yLE', 'xN',
y = df['lookingStraight']]

# Split into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_si

# Define the parameter grid for GradientBoostingClassifier
param_grid_gbc = {
    'learning_rate': [0.01, 0.1],
    'n_estimators': [100, 200],
    'max_depth': [3, 5]
}

# Initialize the GradientBoostingClassifier
gbc = GradientBoostingClassifier()

# Perform grid search with cross-validation
grid_search_gbc = RandomizedSearchCV(estimator=gbc, param_distrib
grid_search_gbc.fit(X_train, y_train)

# Best parameters and best score for GradientBoostingClassifier
best_params_gbc = grid_search_gbc.best_params_
best_score_gbc = grid_search_gbc.best_score_

# Define the parameter grid for SVC
param_distributions_svc = {
    'C': [0.1, 1],
    'kernel': ['linear', 'rbf'],
    'gamma': ['scale', 'auto']
}

# Initialize the SVC
svc = SVC()
```

```

# Use RandomizedSearchCV for SVC
random_search_svc = RandomizedSearchCV(estimator=svc, param_distr
random_search_svc.fit(X_train, y_train)

# Best parameters and best score for SVC
best_params_svc = random_search_svc.best_params_
best_score_svc = random_search_svc.best_score_

# Train the best GradientBoostingClassifier with the found paramet
best_gbc = GradientBoostingClassifier(**best_params_gbc)
best_gbc.fit(X_train, y_train)
y_pred_gbc = best_gbc.predict(X_test)

# Train the best SVC with the found parameters
best_svc = SVC(**best_params_svc)
best_svc.fit(X_train, y_train)
y_pred_svc = best_svc.predict(X_test)

# Calculate F1 Scores for both models
f1_gbc = f1_score(y_test, y_pred_gbc)
f1_svc = f1_score(y_test, y_pred_svc)

print("Best parameters for GradientBoostingClassifier:", best_par
print("Best accuracy score for GradientBoostingClassifier:", best
print("F1 Score for GradientBoostingClassifier:", f1_gbc)

print("Best parameters for SVC:", best_params_svc)
print("Best accuracy score for SVC:", best_score_svc)
print("F1 Score for SVC:", f1_svc)

# Feature importance for GradientBoostingClassifier
feature_importances = best_gbc.feature_importances_
features = X.columns

# Create a DataFrame for feature importances
feature_importances_df = pd.DataFrame({'Feature': features, 'Impo
print("Feature importances for GradientBoostingClassifier:\n", fe

# Misclassification analysis for GradientBoostingClassifier
misclassified_gbc = X_test[(y_test != y_pred_gbc)]
correct_gbc = X_test[(y_test == y_pred_gbc)]

# Summary of misclassified instances
misclassified_summary_gbc = misclassified_gbc.describe()
correct_summary_gbc = correct_gbc.describe()

# Misclassification analysis for SVC
misclassified_svc = X_test[(y_test != y_pred_svc)]
correct_svc = X_test[(y_test == y_pred_svc)]

```

```
# Summary of misclassified instances
```

```
misclassified_summary_svc = misclassified_svc.describe()
```

```
correct_summary_svc = correct_svc.describe()
```

```
print("Summary of misclassified instances for GradientBoostingCla
```

```
print("Summary of correctly classified instances for GradientBoos
```

```
print("Summary of misclassified instances for SVC:\n", misclassif
```

```
print("Summary of correctly classified instances for SVC:\n", cor
```

```
/Users/adrianchavezloya/anaconda3/lib/python3.11/site-packages/sk  
learn/model_selection/_search.py:307: UserWarning: The total spac  
e of parameters 8 is smaller than n_iter=10. Running 8 iteration  
s. For exhaustive searches, use GridSearchCV.
```

```
warnings.warn(
```

```
/Users/adrianchavezloya/anaconda3/lib/python3.11/site-packages/sk  
learn/model_selection/_search.py:307: UserWarning: The total spac  
e of parameters 8 is smaller than n_iter=10. Running 8 iteration  
s. For exhaustive searches, use GridSearchCV.
```

```
warnings.warn(
```

Best parameters for GradientBoostingClassifier: {'n\_estimators': 200, 'max\_depth': 3, 'learning\_rate': 0.01}  
 Best accuracy score for GradientBoostingClassifier: 0.9463231347289319  
 F1 Score for GradientBoostingClassifier: 0.9688888888888889  
 Best parameters for SVC: {'kernel': 'linear', 'gamma': 'scale', 'C': 0.1}  
 Best accuracy score for SVC: 0.9318176008997265  
 F1 Score for SVC: 0.9596412556053813  
 Feature importances for GradientBoostingClassifier:

	Feature	Importance
8	xN	4.372047e-01
10	xRM	2.984812e-01
2	wF	6.889214e-02
12	xLM	4.925377e-02
1	yF	4.189373e-02
13	yLM	3.576459e-02
6	xLE	3.349510e-02
3	hF	1.373653e-02
4	xRE	8.620396e-03
0	xF	7.890654e-03
9	yN	4.477197e-03
5	yRE	2.559169e-04
7	yLE	3.398371e-05
11	yRM	6.820273e-08

Summary of misclassified instances for GradientBoostingClassifier:

		xF	yF	wF	hF	xR
E	yRE \					
count	7.000000	7.000000	7.000000	7.000000	7.000000	7.000000
mean	303.428571	186.571429	115.857143	129.857143	353.142857	215.142857
std	26.095064	25.277507	8.552360	11.066896	18.595955	25.569699
min	279.000000	154.000000	104.000000	116.000000	331.000000	178.000000
25%	283.500000	169.500000	108.500000	122.500000	341.500000	199.000000
50%	288.000000	187.000000	121.000000	127.000000	342.000000	218.000000
75%	324.000000	203.500000	122.000000	136.500000	369.500000	232.500000
max	342.000000	219.000000	125.000000	148.000000	377.000000	247.000000

		xLE	yLE	xN	yN	xRM
yRM \						
count	7.000000	7.000000	7.000000	7.000000	7.000000	7.000000
	7.000000					



mean	392.714286	216.285714	390.714286	239.000000	362.285714
std	21.731040	24.830856	18.988718	23.958297	13.996598
min	370.000000	180.000000	369.000000	204.000000	341.000000
25%	376.500000	202.000000	378.000000	226.000000	355.500000
50%	380.000000	217.000000	381.000000	239.000000	363.000000
75%	413.000000	232.000000	405.500000	253.500000	369.000000
max	420.000000	249.000000	418.000000	271.000000	383.000000

	xLM	yLM
count	7.000000	7.000000
mean	389.428571	264.857143
std	19.973792	23.926475
min	364.000000	228.000000
25%	376.000000	253.500000
50%	381.000000	266.000000
75%	407.500000	279.000000
max	414.000000	295.000000

Summary of correctly classified instances for GradientBoostingClassifier:

		xF	yF	wF	hF	xR
E	yRE \					
count	115.000000	115.000000	115.000000	115.000000	115.000000	115.000000
mean	300.139130	202.286957	112.486957	128.286957	337.782609	229.339130
std	16.682214	38.964168	7.454382	8.129182	17.801923	37.745195
min	264.000000	148.000000	87.000000	105.000000	287.000000	176.000000
25%	286.000000	171.500000	108.000000	124.000000	325.500000	198.500000
50%	301.000000	200.000000	112.000000	129.000000	337.000000	229.000000
75%	310.000000	223.000000	117.000000	134.000000	349.500000	246.000000
max	348.000000	274.000000	133.000000	150.000000	395.000000	300.000000

	xLE	yLE	xN	yN	xRM
yRM \					
count	115.000000	115.000000	115.000000	115.000000	115.000000
mean	384.756522	230.191304	366.330435	254.200000	344.965217

279.443478					
std	17.280753	39.913480	20.822020	39.318805	16.531698
38.644850					
min	334.000000	171.000000	304.000000	196.000000	305.000000
224.000000					
25%	372.000000	199.000000	352.500000	222.500000	331.000000
246.000000					
50%	387.000000	230.000000	364.000000	252.000000	345.000000
273.000000					
75%	396.000000	247.000000	380.500000	274.000000	355.500000
295.000000					
max	437.000000	304.000000	436.000000	329.000000	393.000000
353.000000					

	xLM	yLM
count	115.000000	115.000000
mean	378.860870	279.800000
std	16.402763	40.225941
min	337.000000	219.000000
25%	367.000000	248.000000
50%	380.000000	275.000000
75%	388.000000	294.500000
max	430.000000	356.000000

Summary of misclassified instances for SVC:

		xF	yF	wF	hF	xR
E	yRE \					
count	9.000000	9.000000	9.000000	9.000000	9.000000	9.000000
9.000000						
mean	314.444444	175.666667	111.333333	135.888889	353.000000	204.444444
std	32.384839	22.005681	7.314369	10.409664	36.383375	23.569637
min	276.000000	154.000000	100.000000	116.000000	298.000000	178.000000
25%	282.000000	155.000000	107.000000	128.000000	324.000000	184.000000
50%	331.000000	176.000000	109.000000	139.000000	376.000000	203.000000
75%	343.000000	181.000000	117.000000	142.000000	380.000000	214.000000
max	348.000000	218.000000	122.000000	150.000000	395.000000	246.000000

	xLE	yLE	xN	yN	xRM
yRM \					
count	9.000000	9.000000	9.000000	9.000000	9.000000
9.000000					
mean	396.333333	203.888889	384.666667	229.777778	358.777778
255.222222					
std	35.067791	23.379716	43.373379	22.055486	31.586302

```

21.347001
min    346.000000    177.000000    313.000000    204.000000    308.000000
229.000000
25%    366.000000    187.000000    353.000000    210.000000    333.000000
237.000000
50%    417.000000    203.000000    405.000000    229.000000    374.000000
256.000000
75%    426.000000    206.000000    418.000000    241.000000    385.000000
269.000000
max    437.000000    246.000000    436.000000    267.000000    393.000000
290.000000

```

```

                xLM                yLM
count    9.000000    9.000000
mean    392.333333    254.666667
std     31.496031    22.000000
min     353.000000    227.000000
25%     361.000000    239.000000
50%     412.000000    259.000000
75%     417.000000    261.000000
max     430.000000    291.000000

```

Summary of correctly classified instances for SVC:

```

                xF                yF                wF                hF                xR
E      yRE \
count  113.000000    113.000000    113.000000    113.000000    113.000000
113.000000
mean    299.203540    203.433628    112.787611    127.778761    337.522124
230.442478
std     15.105675     38.757229     7.561003     7.831637     15.575907
37.510698
min     264.000000    148.000000     87.000000    105.000000    287.000000
176.000000
25%     287.000000    172.000000    108.000000    124.000000    328.000000
201.000000
50%     300.000000    202.000000    112.000000    129.000000    338.000000
229.000000
75%     309.000000    223.000000    117.000000    133.000000    348.000000
247.000000
max     347.000000    274.000000    133.000000    148.000000    376.000000
300.000000

```

```

                xLE                yLE                xN                yN                xRM
yRM \
count  113.000000    113.000000    113.000000    113.000000    113.000000
113.000000
mean    384.327434    231.424779    366.380531    255.203540    344.938053
280.398230
std     15.310735     39.629265     18.340684     39.180024     14.860232
38.533969
min     334.000000    171.000000    304.000000    196.000000    305.000000

```

224.000000					
25%	374.000000	202.000000	353.000000	227.000000	332.000000
246.000000					
50%	385.000000	230.000000	365.000000	252.000000	345.000000
274.000000					
75%	396.000000	249.000000	380.000000	274.000000	356.000000
295.000000					
max	420.000000	304.000000	412.000000	329.000000	383.000000
353.000000					

	xLM	yLM
count	113.000000	113.000000
mean	378.442478	280.876106
std	14.716935	40.064264
min	337.000000	219.000000
25%	368.000000	249.000000
50%	380.000000	276.000000
75%	387.000000	295.000000
max	414.000000	356.000000

# Model Training and Evaluation Summary (with new parameter grid for more efficient and performance

## Dataset Overview

- **Binary Response Variable:** Created from `label`, where `lookingStraight` is 1 if `label` is 2, otherwise 0.
- **Features:** Coordinates and dimensions of facial landmarks (e.g., `xF`, `yF`, `wF`, `hF`, etc.)

## Data Splitting

- **Training Set:** 80%
- **Test Set:** 20%
- **Stratified Split:** Ensures balanced class distribution

## Model Selection and Hyperparameter Tuning

Two machine learning models were evaluated: GradientBoostingClassifier and Support Vector Classifier (SVC). We used GridSearchCV for

hyperparameter tuning to find the best combination of parameters that yield the highest accuracy.

**GradientBoostingClassifier:**

- **Best Parameters:** `{'n_estimators': 200, 'max_depth': 3, 'learning_rate': 0.01}`
- **Best Accuracy Score:** 0.946
- **F1 Score:** 0.969

**Support Vector Classifier (SVC):**

- **Best Parameters:** `{'kernel': 'linear', 'gamma': 'scale', 'C': 0.1}`
- **Best Accuracy Score:** 0.932
- **F1 Score:** 0.960

**Feature Importance for GradientBoostingClassifier**

Feature	Importance
xN	0.437205
xRM	0.298481
wF	0.068892
xLM	0.049254
yF	0.041894
yLM	0.035765
xLE	0.033495
hF	0.013737
xRE	0.008620
xF	0.007891

**Misclassification Analysis**

Performed an analysis of misclassified instances to understand where our models struggled. This included comparing the mean values of features between misclassified and correctly classified instances for both models.

## GradientBoostingClassifier:

- **Misclassified Instances:**

- Higher mean values for  $x_F$ ,  $y_F$ ,  $w_F$ , and  $h_F$  compared to correctly classified instances.
- Misclassification may relate to variations in facial landmark positions and dimensions.

## Support Vector Classifier (SVC):

- **Misclassified Instances:**

- Similar patterns in feature means as observed in the GradientBoostingClassifier.
- $x_F$  and  $y_F$  means significantly differ between misclassified and correctly classified instances.

## Optimization for Speed

To improve the runtime of the models, I adjusted the following:

- **GradientBoostingClassifier:** Reduced the number of estimators and controlled the depth of trees.
- **SVC:** Used a linear kernel and optimized the  $C$  parameter to balance complexity and performance.

## Conclusions

- Both models performed well with high accuracy and F1 scores.
- GradientBoostingClassifier slightly outperformed SVC in terms of accuracy and F1 score.
- Feature importance analysis revealed that  $x_N$  and  $x_{RM}$  were the most significant features.
- Misclassification analysis highlighted key areas for further feature engineering and model improvement.
- Parameter adjustments successfully reduced model training and evaluation times without significantly impacting performance.

## Questions

1. Is accuracy the best metric to use in these tasks or would there have been a better one? Explain.
  - Accuracy is often used as a primary evaluation metric for classification tasks, but its suitability depends on the nature of the problem and the data. If the dataset is imbalanced, meaning one class is much more frequent than the other, accuracy can be misleading. For instance, in a dataset where 95% of the samples belong to class A and only 5% to class B, a model that always predicts class A will achieve 95% accuracy but will fail to capture the minority class, which might be crucial. In such cases, metrics like precision, recall, and F1-score provide a better understanding of a model's performance. Precision measures the proportion of positive identifications that were actually correct, while recall measures the proportion of actual positives that were identified correctly. The F1-score is the harmonic mean of precision and recall, providing a balance between the two. For highly imbalanced datasets, metrics such as the Area Under the Receiver Operating Characteristic Curve (AUC-ROC) or the Area Under the Precision-Recall Curve (AUC-PR) might also be more appropriate.

1. Which model gave the "best" result using the metric you chose above?
  - The model that gave the best result using a more appropriate metric like the F1-score or AUC-ROC would be considered the best model. Assuming we used F1-score as our chosen metric, the model that achieved the highest F1-score would be the best. For example, if Model A had an F1-score of 0.85 and Model B had an F1-score of 0.78, then Model A would be considered the best model. Similarly, if we used AUC-ROC and Model A had an AUC-ROC of 0.92 compared to Model B's AUC-ROC of 0.89, Model A would again be considered superior.

### **3. (Bonus) Any other interesting insights from this model or data?**

1. (Bonus) Any other interesting insights from this model or data?
  - Analyzing the model and data might reveal several interesting insights. For instance, certain features might have a stronger correlation with the

target variable, indicating their importance in predicting outcomes. Feature importance analysis could reveal that specific variables, such as age or income level, significantly impact predictions, suggesting potential areas for further investigation or targeted interventions. Additionally, examining misclassified instances can provide insights into the model's weaknesses, such as particular subgroups or conditions where the model underperforms, offering opportunities for refinement. Understanding these aspects can help in improving the model and making more informed decisions based on its predictions.