

Decision Trees and Random Forests

By Adrian Chavez-Loya

A bank is interested in reaching out to customers directly (via phone) to solicit subscriptions to a new product they are planning to offer. The company has over 45k customers and only a small number of phone agents to contact them so targeting those that are most likely to subscribe will maximize their return on investment.

Prior to contacting them, they have asked their Data Science team to analyze customer characteristics for a similar product campaign they ran in the previous year. Your goal is to construct an effective tree-based model to predict whether a customer will subscribe or not.

Relevant Dataset

bank-full.csv

- Response Variable: y

Source of data:

<https://archive.ics.uci.edu/ml/datasets/Bank+Marketing>

Task 1: Import data and construct a random 70/15/15 train/val/test split. Make sure to dummy code categorical variables.

```
In [ ]: import pandas as pd

df = pd.read_csv('bank-full.csv', delimiter = ";") # Import with
df.head()
```

Out []:

	age	job	marital	education	default	balance	housing	loan	cont
0	58	management	married	tertiary	no	2143	yes	no	unknc
1	44	technician	single	secondary	no	29	yes	no	unknc
2	33	entrepreneur	married	secondary	no	2	yes	yes	unknc
3	47	blue-collar	married	unknown	no	1506	yes	no	unknc
4	33	unknown	single	unknown	no	1	no	no	unknc

In []:

```
# Random 70/15/15 train/validation/test split
import sklearn as skl
from sklearn.model_selection import train_test_split

# Dummy variable categ. variables
df = pd.get_dummies(df, drop_first=True)
print(df.head)
```

	<bound method NDFrame.head of	age	balance	day	duration
campaign	pdays	previous	\		
0	58	2143	5	261	1 -1 0
1	44	29	5	151	1 -1 0
2	33	2	5	76	1 -1 0
3	47	1506	5	92	1 -1 0
4	33	1	5	198	1 -1 0
...
45206	51	825	17	977	3 -1 0
45207	71	1729	17	456	2 -1 0
45208	72	5715	17	1127	5 184 3
45209	57	668	17	508	4 -1 0
45210	37	2971	17	361	2 188 11

	job_blue-collar	job_entrepreneur	job_housemaid	...	mon
th_jun	\				
0	False	False	False	...	
False					
1	False	False	False	...	
False					
2	False	True	False	...	
False					
3	True	False	False	...	
False					
4	False	False	False	...	
False					
...	
...					
45206	False	False	False	...	
False					
45207	False	False	False	...	
False					
45208	False	False	False	...	
False					
45209	True	False	False	...	
False					
45210	False	True	False	...	
False					

	month_mar	month_may	month_nov	month_oct	month_sep	pou
tcome_other	\					
0	False	True	False	False	False	
False						
1	False	True	False	False	False	
False						
2	False	True	False	False	False	
False						
3	False	True	False	False	False	
False						
4	False	True	False	False	False	

```

False
...
...
...
45206      False      False      True      False      False
False
45207      False      False      True      False      False
False
45208      False      False      True      False      False
False
45209      False      False      True      False      False
False
45210      False      False      True      False      False
True

```

```

      poutcome_success  poutcome_unknown  y_yes
0                False                True  False
1                False                True  False
2                False                True  False
3                False                True  False
4                False                True  False
...
45206                False                True   True
45207                False                True   True
45208                 True                False  True
45209                False                True  False
45210                False                False  False

```

```
[45211 rows x 43 columns]>
```

```

In [ ]: # Features and target
X = df.drop('y_yes', axis=1)
y = df['y_yes']

```

```

In [ ]: X_train, X_temp, y_train, y_temp = train_test_split(X, y, test_si
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, t
# Dimensions of train, val., and test set:
print(f"Train set: {X_train.shape}, {y_train.shape}")
print(f"Validation set: {X_val.shape}, {y_val.shape}")
print(f"Test set: {X_test.shape}, {y_test.shape}")

```

```

Train set: (31647, 42), (31647,)
Validation set: (6782, 42), (6782,)
Test set: (6782, 42), (6782,)

```

Task 2: Use Cost-Complexity Pruning to find the optimal depth for a Decision Tree Classifier.

Note: "Optimal" is subjective. Feel free to choose a shallower more interpretable tree or a slightly deeper more accurate tree. Simply explain why you chose the depth you did.


```
In [ ]: from sklearn.tree import DecisionTreeClassifier
        from sklearn.model_selection import cross_val_score
        import matplotlib.pyplot as plt
```

```
In [ ]: clf = DecisionTreeClassifier(random_state=42)
        path = clf.cost_complexity_pruning_path(X_train, y_train)
        ccp_alphas = path.ccp_alphas

        # Cross-validation to find optimal alpha
        clf_scores = []
        for ccp_alpha in ccp_alphas:
            clf = DecisionTreeClassifier(random_state=42, ccp_alpha=ccp_alpha)
            scores = cross_val_score(clf, X_train, y_train, cv=5)
            clf_scores.append(scores.mean())
```

```
In [ ]: plt.figure()
        plt.plot(ccp_alphas, clf_scores, marker='o', drawstyle='steps-post')
        plt.xlabel('Alpha')
        plt.ylabel('Cross-Validated Accuracy')
        plt.title('Alpha vs. Cross-Validated Accuracy')
        plt.show()

        # get optimal alpha value
        optimal_alpha = ccp_alphas[clf_scores.index(max(clf_scores))]
        clf_optimal = DecisionTreeClassifier(random_state=42, ccp_alpha=optimal_alpha)
        clf_optimal.fit(X_train, y_train)
        print(f"Optimal alpha: {optimal_alpha}")
        print(f"Optimal tree depth: {clf_optimal.get_depth()}")
```



Optimal alpha: 0.000317039596610124

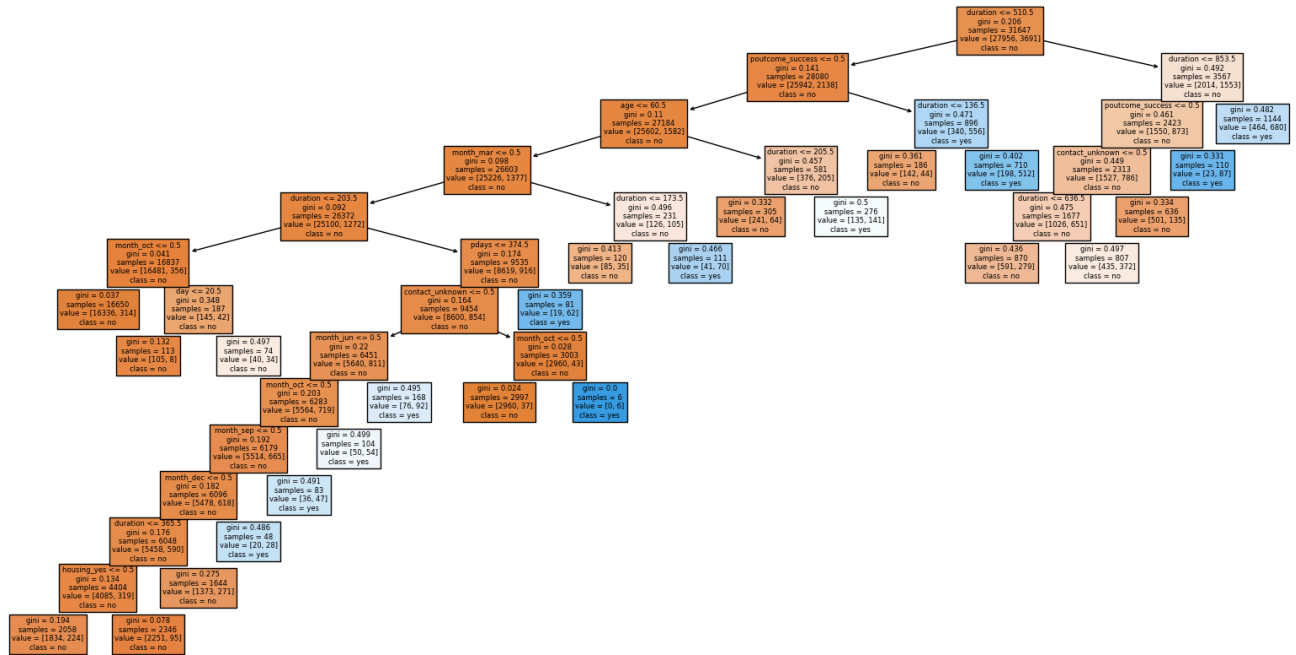
Optimal tree depth: 13

Task 3: Visualize the "Optimal" tree you fit and interpret the first few splits to the best of your ability.

```
In [ ]: from sklearn.tree import plot_tree

# Convert feature names to a list
feature_names = X.columns.tolist()

# Plot the optimal decision tree
plt.figure(figsize=(20,10))
plot_tree(clf_optimal, filled=True, feature_names=feature_names,
plt.show())
```



The decision tree model starts by splitting on `duration <= 203.5`, with a Gini index of 0.092 and a majority class of `no`. For samples with `duration <= 203.5`, the next split is on `month_oct <= 0.5`, achieving a Gini index of 0.041, still dominated by the `no` class. If `month_oct > 0.5`, it further splits on `day <= 20.5` with a moderate Gini index of 0.348. For `duration > 203.5`, the model splits again on `duration <= 510.5` with a Gini index of 0.206, maintaining `no` as the majority class. Subsequent splits for `duration <= 203.5` and `month_oct <= 0.5` include `month_jun <= 0.5`, achieving a Gini index of 0.022. Overall, `duration` and `month` features are critical in classifying the samples.

Task 4: Calculate the k-fold CV accuracy for predicting 'yes' or 'no' using a Decision Tree Classifier.

```
In [ ]: from sklearn.model_selection import cross_val_score

# Using the optimal Decision Tree from Task 2
cv_scores = cross_val_score(clf_optimal, X, y, cv=5)
print(f"5-Fold CV accuracy: {cv_scores.mean()}")
```

5-Fold CV accuracy: 0.6643467570446618

- 5-fold CV accuracy is 66.4 percent

Task 5: Fit a Random Forest Classifier and use the OOB accuracy to choose the optimal number of trees and/or variables sampled at each split.

Note: The argument for number of trees is `n_estimators` in scikit-learn and the number of variables to consider at each split is `max_features`.

```
In [ ]: from sklearn.ensemble import RandomForestClassifier

# Fit a Random Forest with default parameters to check OOB score
rf_clf = RandomForestClassifier(oob_score=True, random_state=42)
rf_clf.fit(X_train, y_train)
print(f"OOB score: {rf_clf.oob_score_}")

# Experiment with different number of trees and max_features to o
best_oob_score = 0
best_params = {}

for n_estimators in [50, 100, 200]:
    for max_features in [None, 'sqrt', 'log2']:
        rf_clf = RandomForestClassifier(n_estimators=n_estimators,
                                       max_features=max_features,
                                       random_state=42)
        rf_clf.fit(X_train, y_train)
        if rf_clf.oob_score_ > best_oob_score:
            best_oob_score = rf_clf.oob_score_
            best_params = {'n_estimators': n_estimators, 'max_fea

print(f"Best OOB score: {best_oob_score}")
print(f"Best parameters: {best_params}")

# Fit the Random Forest with the best parameters
rf_best = RandomForestClassifier(n_estimators=best_params['n_esti
rf_best.fit(X_train, y_train)
print(f"Best OOB Score: {rf_best.oob_score_}")
print(f"Best Parameters: {best_params}")
```

OOB score: 0.901602047587449

Best OOB score: 0.9044143204727146

Best parameters: {'n_estimators': 200, 'max_features': 'log2'}

Best OOB Score: 0.9044143204727146

Best Parameters: {'n_estimators': 200, 'max_features': 'log2'}

Task 6: Report impurity based feature importance for your final selected model in a


bar chart.

Optional: If you are curious, try running permutation importance (note: this may take a while) and compare the two

```
In [ ]: import numpy as np

importances = rf_best.feature_importances_
indices = np.argsort(importances)[::-1]
features = X_train.columns

# Order highest feature importances
plt.figure(figsize=(15, 8))
plt.title("Feature Importances")
plt.bar(range(X_train.shape[1]), importances[indices], align="center")
plt.xticks(range(X_train.shape[1]), features[indices], rotation=90)
plt.xlim([-1, X_train.shape[1]])
plt.show()
```



Task 7: Choose the model with the best cross-validated or OOB accuracy between Decision Trees and Random Forests and provide the

Partial Dependence Plots for the 5 most important variables.

```
In [ ]: from sklearn.inspection import PartialDependenceDisplay

top_5_features = features[indices[:5]] # Top 5 Features
fig, ax = plt.subplots(figsize=(15, 10))
PartialDependenceDisplay.from_estimator(rf_best, X_train, top_5_f
plt.suptitle('Partial Dependence Plots')
plt.subplots_adjust(top=0.9) # Adjust title position
plt.show()
```

