

Generalized Linear Models (GLMs)

By Adrian Chavez-Loya

Background

In 2014, the analytics website, FiveThirtyEight, published an analysis of the gender disparity in appearances in comic books at the two large studios, Marvel and DC. You're working for a small upstart comic book company and you want to work to understand what characteristics are associated with a reduction in appearances and thereby learn where gaps in representation can be filled by your firm. Using the data collected by FiveThirtyEight from the Marvel and DC wiki pages, you will fit an appropriate GLM to predict the number of appearances given this data.

Relevant Datasets

- `marvel-wikia-data.csv`
- `dc-wikia-data.csv`

Article: <https://fivethirtyeight.com/features/women-in-comic-books/>

Aggregated Data: <https://github.com/fivethirtyeight/data/tree/master/comic-characters>

Data Sources:

- http://dc.wikia.com/wiki/Main_Page
- http://marvel.wikia.com/Main_Page

Task 1: Concatenate the Marvel and DC datasets into a single dataset.

The function `pd.concat` will be useful for this. You should also create a new column in each dataset before concatenating that represents the

studio for each. This will allow us to distinguish between the impact of one studio vs the other.

```
In [ ]: # Imported data and data types to perform data cleaning
import pandas as pd
import numpy as np
marvel_data = pd.read_csv('marvel-wikia-data.csv')
dc_data = pd.read_csv('dc-wikia-data.csv')
print("Marvel Data Types and Head:")
print(marvel_data.dtypes)
print(marvel_data.head())
print("\nDC Data Types and Head:")
print(dc_data.dtypes)
print(dc_data.head())
```

Marvel Data Types and Head:

page_id	int64
name	object
urlslug	object
ID	object
ALIGN	object
EYE	object
HAIR	object
SEX	object
GSM	object
ALIVE	object
APPEARANCES	float64
FIRST APPEARANCE	object
Year	float64

dtype: object

	page_id	name \
0	1678	Spider-Man (Peter Parker)
1	7139	Captain America (Steven Rogers)
2	64786	Wolverine (James \"Logan\" Howlett)
3	1868	Iron Man (Anthony \"Tony\" Stark)
4	2460	Thor (Thor Odinson)

	urlslug	ID \
0	\\/Spider-Man_(Peter_Parker)	Secret Identity
1	\\/Captain_America_(Steven_Rogers)	Public Identity
2	\\/Wolverine_(James_%22Logan%22_Howlett)	Public Identity
3	\\/Iron_Man_(Anthony_%22Tony%22_Stark)	Public Identity
4	\\/Thor_(Thor_Odinson)	No Dual Identity

	ALIGN	EYE	HAIR	SEX	G
SM \					
0	Good Characters	Hazel Eyes	Brown Hair	Male Characters	N
aN					
1	Good Characters	Blue Eyes	White Hair	Male Characters	N
aN					
2	Neutral Characters	Blue Eyes	Black Hair	Male Characters	N
aN					
3	Good Characters	Blue Eyes	Black Hair	Male Characters	N
aN					
4	Good Characters	Blue Eyes	Blond Hair	Male Characters	N
aN					

	ALIVE	APPEARANCES	FIRST APPEARANCE	Year
0	Living Characters	4043.0	Aug-62	1962.0
1	Living Characters	3360.0	Mar-41	1941.0
2	Living Characters	3061.0	Oct-74	1974.0
3	Living Characters	2961.0	Mar-63	1963.0
4	Living Characters	2258.0	Nov-50	1950.0

DC Data Types and Head:

```

page_id      int64
name         object
urlslug      object
ID           object
ALIGN        object
EYE          object
HAIR         object
SEX          object
GSM          object
ALIVE        object
APPEARANCES  float64
FIRST APPEARANCE object
YEAR         float64
dtype: object

```

```

      page_id      name
urlslug \
0      1422      Batman (Bruce Wayne)      \wiki\Batman_(B
ruce_Wayne)
1      23387      Superman (Clark Kent)      \wiki\Superman_
(Clark_Kent)
2      1458      Green Lantern (Hal Jordan)      \wiki\Green_Lantern_
(Hal_Jordan)
3      1659      James Gordon (New Earth)      \wiki\James_Gordon_
(New_Earth)
4      1576      Richard Grayson (New Earth)      \wiki\Richard_Grayson_
(New_Earth)

```

```

      ID      ALIGN      EYE      HAIR
SEX \
0 Secret Identity Good Characters Blue Eyes Black Hair Male
Characters
1 Secret Identity Good Characters Blue Eyes Black Hair Male
Characters
2 Secret Identity Good Characters Brown Eyes Brown Hair Male
Characters
3 Public Identity Good Characters Brown Eyes White Hair Male
Characters
4 Secret Identity Good Characters Blue Eyes Black Hair Male
Characters

```

```

      GSM      ALIVE      APPEARANCES      FIRST APPEARANCE      YEAR
0 NaN Living Characters      3093.0      1939, May      1939.0
1 NaN Living Characters      2496.0      1986, October      1986.0
2 NaN Living Characters      1565.0      1959, October      1959.0
3 NaN Living Characters      1316.0      1987, February      1987.0
4 NaN Living Characters      1237.0      1940, April      1940.0

```

```

In [ ]: marvel_data['Studio'] = 'Marvel' # Labeled by studio
dc_data['Studio'] = 'DC'

```

```
# Combined 2 data sets
combined_data = pd.concat([marvel_data, dc_data], ignore_index=True)
print("Combined Data Types and Head:")
print(combined_data.dtypes)
print(combined_data.head())
```

Combined Data Types and Head:

page_id	int64
name	object
urlslug	object
ID	object
ALIGN	object
EYE	object
HAIR	object
SEX	object
GSM	object
ALIVE	object
APPEARANCES	float64
FIRST APPEARANCE	object
Year	float64
Studio	object
YEAR	float64

dtype: object

	page_id	name \
0	1678	Spider-Man (Peter Parker)
1	7139	Captain America (Steven Rogers)
2	64786	Wolverine (James \"Logan\" Howlett)
3	1868	Iron Man (Anthony \"Tony\" Stark)
4	2460	Thor (Thor Odinson)

	urlslug	ID \
0	\\/Spider-Man_(Peter_Parker)	Secret Identity
1	\\/Captain_America_(Steven_Rogers)	Public Identity
2	\\/Wolverine_(James_%22Logan%22_Howlett)	Public Identity
3	\\/Iron_Man_(Anthony_%22Tony%22_Stark)	Public Identity
4	\\/Thor_(Thor_Odinson)	No Dual Identity

	ALIGN	EYE	HAIR	SEX	G
SM \					
0	Good Characters	Hazel Eyes	Brown Hair	Male Characters	N
aN					
1	Good Characters	Blue Eyes	White Hair	Male Characters	N
aN					
2	Neutral Characters	Blue Eyes	Black Hair	Male Characters	N
aN					
3	Good Characters	Blue Eyes	Black Hair	Male Characters	N
aN					
4	Good Characters	Blue Eyes	Blond Hair	Male Characters	N
aN					

	ALIVE	APPEARANCES	FIRST APPEARANCE	Year	Studi
o	YEAR				
0	Living Characters	4043.0	Aug-62	1962.0	Marve
l	NaN				
1	Living Characters	3360.0	Mar-41	1941.0	Marve
l	NaN				

2	Living Characters	3061.0	Oct-74	1974.0	Marvel
1	NaN				
3	Living Characters	2961.0	Mar-63	1963.0	Marvel
1	NaN				
4	Living Characters	2258.0	Nov-50	1950.0	Marvel
1	NaN				

```
In [ ]: # Subsetting
relevant_columns = ['ALIGN', 'SEX', 'ALIVE', 'APPEARANCES', 'ID',
subset_data = combined_data[relevant_columns].copy()
subset_data.loc[:, 'SEX'] = subset_data['SEX'].astype(str) # Sex
subset_data.loc[:, 'SEX'].replace('nan', 'Unknown', inplace=True)
subset_data.loc[:, 'Is_Male'] = subset_data['SEX'].apply(lambda x
filtered_data = subset_data[subset_data['APPEARANCES'] > 1] # Rem
print("Filtered Data Types and Head:")
print(filtered_data.dtypes)
print(filtered_data.head())
```

Filtered Data Types and Head:

```
ALIGN      object
SEX         object
ALIVE      object
APPEARANCES float64
ID          object
Studio     object
Is_Male    int64
dtype: object
```

		ALIGN	SEX	ALIVE	APPEAR
ANCES \					
0	Good Characters	Male Characters	Living Characters		4
043.0					
1	Good Characters	Male Characters	Living Characters		3
360.0					
2	Neutral Characters	Male Characters	Living Characters		3
061.0					
3	Good Characters	Male Characters	Living Characters		2
961.0					
4	Good Characters	Male Characters	Living Characters		2
258.0					

	ID	Studio	Is_Male
0	Secret Identity	Marvel	1
1	Public Identity	Marvel	1
2	Public Identity	Marvel	1
3	Public Identity	Marvel	1
4	No Dual Identity	Marvel	1

Task 2: Subset the data to relevant variables and observations.

If you want to limit the number of levels, a good list of variables would be: `ALIGN`, `SEX`, `ALIVE`, `APPEARANCES`, `ID`, `Studio`. Given that this dataset includes a few `SEX` categories with very few observations, create a new binary variable for a character's `SEX` being Male or not Male. Also, remove any characters that only appear once.

```
In [ ]: # Final data cleaning
filtered_data = filtered_data.dropna(subset=['ALIGN', 'SEX', 'ALI
filtered_data['ALIGN'] = filtered_data['ALIGN'].astype('category')
filtered_data['SEX'] = filtered_data['SEX'].astype('category')
filtered_data['ALIVE'] = filtered_data['ALIVE'].astype('category')
filtered_data['Studio'] = filtered_data['Studio'].astype('categor
print("Cleaned Data Types and Head:")
print(filtered_data.dtypes)
print(filtered_data.head())
```


Filtered Data Types and Head:
ALIGN object
SEX object
ALIVE object
APPEARANCES float64
ID object
Studio object
Is_Male int64
dtype: object

		ALIGN	SEX	ALIVE	APPEAR
ANCES \					
0	Good Characters	Male Characters	Living Characters		4
043.0					
1	Good Characters	Male Characters	Living Characters		3
360.0					
2	Neutral Characters	Male Characters	Living Characters		3
061.0					
3	Good Characters	Male Characters	Living Characters		2
961.0					
4	Good Characters	Male Characters	Living Characters		2
258.0					

	ID	Studio	Is_Male
0	Secret Identity	Marvel	1
1	Public Identity	Marvel	1
2	Public Identity	Marvel	1
3	Public Identity	Marvel	1
4	No Dual Identity	Marvel	1

Cleaned Data Types and Head:
ALIGN category
SEX category
ALIVE category
APPEARANCES float64
ID object
Studio category
Is_Male int64
dtype: object

		ALIGN	SEX	ALIVE	APPEAR
ANCES \					
0	Good Characters	Male Characters	Living Characters		4
043.0					
1	Good Characters	Male Characters	Living Characters		3
360.0					
2	Neutral Characters	Male Characters	Living Characters		3
061.0					
3	Good Characters	Male Characters	Living Characters		2
961.0					
4	Good Characters	Male Characters	Living Characters		2
258.0					

		ID	Studio	Is_Male
0	Secret Identity	Marvel	1	
1	Public Identity	Marvel	1	
2	Public Identity	Marvel	1	
3	Public Identity	Marvel	1	
4	No Dual Identity	Marvel	1	

Task 3: Split your data into train/test and fit an appropriate GLM to the training data.

It will be up to you to determine the appropriate choice of distribution or family of the GLM. Look at residual plots and see if there are any red flags with this model.

```
In [ ]: import statsmodels.api as sm
from sklearn.model_selection import train_test_split
relevant_data = combined_data[['ALIGN', 'SEX', 'ALIVE', 'APPEARAN

relevant_data['SEX'] = relevant_data['SEX'].fillna('Unknown') #
relevant_data['Is_Male'] = relevant_data['SEX'].apply(lambda x: 1
filtered_data = relevant_data[relevant_data['APPEARANCES'] > 1]
filtered_data = filtered_data.drop(columns=['SEX']) # Dropped sex
X = pd.get_dummies(filtered_data[['ALIGN', 'ALIVE', 'Is_Male', 'S
y = filtered_data['APPEARANCES']

# Split into train/test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_si
X_train_sm = sm.add_constant(X_train)
```

```

/var/folders/vq/l_8lvyx12cxb7kcq563rstwh0000gn/T/ipykernel_4731/4
283741262.py:11: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFram
e.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/p
andas-docs/stable/user_guide/indexing.html#returning-a-view-versu
s-a-copy
    relevant_data['SEX'] = relevant_data['SEX'].fillna('Unknown')
# Replace NaN with 'Unknown'
/var/folders/vq/l_8lvyx12cxb7kcq563rstwh0000gn/T/ipykernel_4731/4
283741262.py:14: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFram
e.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/p
andas-docs/stable/user_guide/indexing.html#returning-a-view-versu
s-a-copy
    relevant_data['Is_Male'] = relevant_data['SEX'].apply(lambda x:
1 if 'Male' in x else 0)

```

```

In [ ]: # Convert boolean columns to integers
X_train = X_train.astype(int)
print("\nX_train Numpy Array Type after Conversion:")
print(np.asarray(X_train).dtype)

```

```

X_train Numpy Array Type after Conversion:
int64

```

```

In [ ]: # Multicollinearity check with VIF
from statsmodels.stats.outliers_influence import variance_inflati
import pandas as pd
X_train_with_const = sm.add_constant(X_train)
vif = pd.DataFrame()
vif['Feature'] = X_train_with_const.columns
vif['VIF'] = [variance_inflation_factor(X_train_with_const.values
print(vif)

```

	Feature	VIF
0	const	9.822023
1	Is_Male	1.015725
2	ALIGN_Good Characters	1.107689
3	ALIGN_Neutral Characters	1.091771
4	ALIGN_Reformed Criminals	1.000572
5	ALIVE_Living Characters	1.008271
6	Studio_Marvel	1.016270

```

In [ ]: # Fit Poisson GLM
import statsmodels.api as sm

```

```
poisson_model = sm.GLM(y_train, X_train, family=sm.families.Poiss
print("Poisson GLM Summary:")
print(poisson_model.summary())
```

Poisson GLM Summary:

Generalized Linear Model Regression Results

```
=====
=====
Dep. Variable:                APPEARANCES    No. Observations:
12807
Model:                        GLM            Df Residuals:
12801
Model Family:                Poisson        Df Model:
5
Link Function:               Log           Scale:
1.0000
Method:                      IRLS          Log-Likelihood:
-5.2508e+05
Date:                        Fri, 26 Jul 2024    Deviance:
9.9967e+05
Time:                        10:52:08          Pearson chi2:
4.39e+06
No. Iterations:              7              Pseudo R-squ. (CS):
0.1889
Covariance Type:            nonrobust
=====
=====
```

			coef	std err	z	P>
z	[0.025	0.975]				
Is_Male			0.8131	0.004	191.990	0.
000	0.805	0.821				
ALIGN_Good Characters			1.8575	0.005	406.103	0.
000	1.849	1.866				
ALIGN_Neutral Characters			1.3113	0.006	202.895	0.
000	1.299	1.324				
ALIGN_Reformed Criminals			-0.0435	0.302	-0.144	0.
885	-0.634	0.548				
ALIVE_Living Characters			1.2612	0.005	243.174	0.
000	1.251	1.271				
Studio_Marvel			0.3450	0.004	88.183	0.
000	0.337	0.353				

```
=====
=====
```

Task 4: Interpret coefficients in the context of our original research question of "what

characteristics are associated with a reduction in appearances?"

Note: You can ignore any broken assumptions at this point and simply treat it as an effective model.

Interpretation of Coefficients in the Context of the Research Question

The original research question is: "What characteristics are associated with a reduction in appearances?"

Given the Poisson GLM summary, we interpret the coefficients to understand how each characteristic influences the number of appearances.

Summary of Coefficients

1. **Is_Male (Coefficient: 0.8131)**

- Being male is associated with an increase in the number of appearances. Male characters have a higher expected number of appearances compared to non-male characters.

2. **ALIGN_Good Characters (Coefficient: 1.8575)**

- Being classified as a good character is strongly associated with an increase in appearances. Good characters are expected to appear significantly more often than non-good characters.

3. **ALIGN_Neutral Characters (Coefficient: 1.3113)**

- Neutral characters see an increase in the number of appearances. Neutral characters are expected to appear more frequently compared to non-neutral characters.

4. **ALIGN_Reformed Criminals (Coefficient: -0.0435)**

- Being a reformed criminal is associated with a slight reduction in appearances. Although this coefficient is not statistically significant ($p = 0.885$), it suggests a potential tendency for reformed criminals to have fewer appearances.

5. **ALIVE_Living Characters (Coefficient: 1.2612)**

- Living characters are associated with an increase in appearances. Characters who are alive tend to appear more frequently than those who are not.

6. **Studio_Marvel (Coefficient: 0.3450)**

- Characters associated with Marvel Studios are more likely to have a higher number of appearances. Being part of the Marvel universe contributes to more frequent appearances.

Conclusion

- **Characteristics Associated with a Reduction in Appearances:**
 - **ALIGN_Reformed Criminals:** Despite being the only negative coefficient, it is not statistically significant. It suggests a potential, but weak, reduction in appearances.
- **Overall Finding:**
 - There are no strong characteristics identified that significantly reduce the number of appearances. Most characteristics analyzed, including being male, being a good or neutral character, being alive, and being part of Marvel Studios, are associated with an increase in appearances.